

Prompt Engineering Principles

Building Effective Prompts Through Iteration

Franck Albinet

franckalbinet@gmail.com

Independent Data Science & AI Consultant

October 8, 2025

Our Goal for Today

Building a Paper Extraction Prompt

By the end of today, you'll have:

- A systematic understanding of prompt engineering principles
- A working prompt for extracting paper information
- Experience with iterative refinement
- Confidence to tackle other prompting tasks

Key insight: We'll build ONE prompt together, adding complexity step by step

The Seven Building Blocks

Prompting Fundamentals

The Components of Effective Prompts

A well-structured prompt typically includes:

1. **Role and Objective** - Define the AI's persona and goal
2. **Instructions / Response Rules** - Clear, specific directives
3. **Context** - The relevant background information
4. **Examples** - Few-shot prompting for guidance
5. **Reasoning Steps** - Chain-of-thought prompting
6. **Output Formatting** - Structure and constraints
7. **Delimiters** - Clear organization of prompt sections

Reference: See the full guidelines in [resources/prompt_fundamentals.md](#)

Why These Building Blocks?

Addressing the Gulf of Specification

Remember the Three Gulfs from Day 1?

Gulf of Specification: The gap between what you want and what you communicate

These building blocks help you:

- Make your intent explicit
- Reduce ambiguity
- Get consistent outputs
- Bridge that gulf systematically

Our Iterative Approach

How We'll Work Today

We'll build a paper extraction prompt in cycles:

1. **Start simple** - Basic instruction only
2. **Add complexity** - One building block at a time
3. **Test and observe** - What changed?
4. **Refine** - Based on what we learned

Important: This is how real prompt engineering works - iteration, not perfection on the first try

Cycle 1: Basic Instruction

Building Block: Instructions

Starting at the Simplest Level

The core component: Clear, specific directives for the task

For paper extraction, a basic instruction might be:

“Summarize this scientific paper.”

or

“Extract key information from this paper.”

Question: What’s missing from these instructions?

What Makes Good Instructions?

Key Characteristics

- **Specific:** Define exactly what to do
- **Unambiguous:** One clear interpretation
- **Actionable:** The AI knows what action to take
- **Complete:** Covers the full scope of the task

But even good instructions alone aren't enough...

Hands-On: Cycle 1

Your Turn - Basic Instruction

Task: Create a basic instruction-only prompt for paper extraction

See: [session-2c-hands-on.md](#) for detailed guidance

Time: 15 minutes

Remember to note:

- What output did you get?
- What was good?
- What was missing or inconsistent?

Debrief: Cycle 1

What Did We Learn?

Group discussion:

- What worked with basic instructions?
- What problems did you encounter?
- What would you want to improve?

Common pattern: The AI gives you *something*, but probably not quite what you wanted

Cycle 2: Adding Role & Objective

Building Blocks: Role and Objective

Setting the Context

Role: Define the AI's persona

- Helps the AI “understand” its perspective
- Influences tone and approach

Objective: State the overall goal

- Provides purpose beyond the immediate instruction
- Helps prioritize what matters

Example: Role & Objective

For Paper Extraction

Role:

“You are a scientific literature extraction specialist with expertise in [domain] research.”

Objective:

“Generate a comprehensive, structured summary of the provided scientific paper optimized for literature review preparation.”

Notice: This adds context *around* your instructions

Why This Matters

Influencing AI Behavior

Adding role and objective:

- Activates relevant “knowledge” in the model
- Sets appropriate tone and detail level
- Provides decision-making criteria
- Makes the task feel more bounded

Key point: You're not just saying WHAT to do, but WHO is doing it and WHY

Hands-On: Cycle 2

Your Turn - Add Role & Objective

Task: Enhance your Cycle 1 prompt by adding: - An appropriate role - A clear objective

See: [session-2c-hands-on.md](#) for guidance

Time: 15 minutes

Compare: How does the output differ from Cycle 1?

Debrief: Cycle 2

What Changed?

Group discussion:

- How did adding role/objective affect the output?
- Was it more aligned with what you wanted?
- What issues persist?

Next challenge: The output might be better, but is it *structured* the way you need?

Cycle 3: Output Formatting

Building Block: Output Formatting

Defining Structure and Constraints

The challenge: Even good content isn't useful if it's not structured properly

Output formatting specifies:

- Structure (sections, lists, tables)
- Format (JSON, Markdown, plain text)
- Constraints (length, level of detail)
- Required elements

Example: Output Formatting

For Paper Extraction

****Format required:****

- Title and citation information
- Abstract (original, verbatim)
- Key takeaways by section (bullet points)
- Methodology summary
- Main results with statistics
- Figures and tables described
- Maximum 2000 tokens

Critical: Be explicit about structure - don't assume the AI will guess correctly

Why Formatting Matters

Consistency and Usability

Explicit formatting requirements:

- Ensure consistency across multiple papers
- Make outputs machine-readable (if needed)
- Enable comparison and synthesis later
- Reduce variability in results

This directly addresses: Gulf of Specification AND Gulf of Generalization

Hands-On: Cycle 3

Your Turn - Add Output Formatting

Task: Add detailed formatting requirements to your prompt

See: [session-2c-hands-on.md](#) for guidance

Time: 20 minutes

Pay attention to: How much more consistent is the output now?

Debrief: Cycle 3

Progress Check

Group discussion:

- Is the structure now what you need?
- Are outputs more consistent?
- What would still improve it?

Pattern emerging: More specific prompts → more consistent outputs

Cycle 4: Context & Examples

Building Blocks: Context and Examples

Providing Background and Guidance

Context: Relevant background information

- Domain-specific knowledge
- Constraints or requirements
- What the output will be used for

Examples (Few-shot): Sample input-output pairs

- Shows exactly what you want
- Clarifies ambiguous instructions
- Demonstrates edge cases

When to Use Examples

The Power of Few-Shot Prompting

Examples are especially valuable when:

- Instructions are complex or nuanced
- You want a specific style or format
- Edge cases need to be handled
- The task has implicit requirements

Trade-off: Examples take up token space but can dramatically improve results

Example: Adding Context

For Paper Extraction

Context example:

****Context:**** These summaries will be used as input for literature review synthesis. Focus on:

- Quantitative results with statistical measures
- Methodological approaches that could be compared
- Key findings relevant to [specific research question]

****Exclude:**** Author affiliations, acknowledgments, funding information

Hands-On: Cycle 4

Your Turn - Add Context

Task: Add relevant context to your prompt

- What will the summary be used for?
- What should be prioritized?
- What should be excluded?

Optional: Add an example if helpful

See: [session-2c-hands-on.md](#) for guidance

Time: 20 minutes

Debrief: Cycle 4

Refinement Continues

Group discussion:

- Did context help focus the output?
- If you used examples, how did they help?
- Is the output now “good enough” for your needs?

Key question: When is a prompt “good enough”? (It depends on your use case!)

Advanced Building Blocks

Reasoning Steps

Chain-of-Thought Prompting

When to use: - Complex analysis tasks - Multi-step reasoning needed - Want to see the AI's thinking process

Example: "Think step by step before answering"

For paper extraction: Less critical, but useful for evaluation tasks

Delimiters

Organizing Complex Prompts

Purpose: Clear separation of prompt sections

Common markers: ###, ---, XML tags

For paper extraction: Helpful as prompts grow more complex

When to Use Reasoning Steps

Chain-of-Thought Prompting

Best for tasks requiring:

- Multi-step analysis
- Comparison or evaluation, complex decision-making, verification of results

Example for evaluation task:

“Before rating this paper’s methodology, first identify the study design, then assess each component (sampling, measurement, analysis), and finally provide an overall rating with justification.”

Cycle 5 (Optional): Advanced Features

Your Choice

If time allows and interest exists:

Add reasoning steps or better organization with delimiters

See: [session-2c-hands-on.md](#) for guidance






Time: 15-20 minutes

The Complete Picture

From Simple to Sophisticated

Your Prompt Evolution

You've built a prompt from:

1.  Basic instruction
2.  + Role & Objective
3.  + Output Formatting
4.  + Context (& Examples)
5.  (Optional) Reasoning Steps & Delimiters

This is the iterative process of real prompt engineering

The Professional Version

Where This Can Lead

Example: The `prompt_paper_representation.md` shows a fully developed version

- All building blocks integrated
- Comprehensive specifications
- Error handling
- Quality standards

Important truth: This wasn't written perfectly from the start - it evolved through iteration

Honesty Moment

How the Example Prompt Was Created

The truth about `prompt_paper_representation.md`:

1. I built an initial version using these principles
2. I tested it on several papers
3. I identified failure modes
4. **I asked Claude AI to refine it** based on my observations
5. I tested again and refined further

Lesson: Use AI to help refine prompts, but start with human-led iteration first

Why Start Manual?

The Value of Doing It Yourself First

Manual iteration first teaches you:

- What you actually want to achieve
- Hidden requirements and edge cases
- How to recognize failure modes
- How to evaluate quality

Then AI refinement helps with:

- Phrasing and clarity
- Comprehensive coverage
- Formatting and organization

Group Comparison

Sharing Our Work

What Did We Build?

Activity: Let's see a few examples

- Volunteers share their final prompts
- What building blocks did you emphasize?
- What worked well for your domain?
- What challenges remain?

Time: 15-20 minutes

Common Patterns

What We Typically See

Across different domains and approaches:

- **Variation in detail level** - Some want more, some less
- **Domain-specific requirements** - Different fields need different information
- **Trade-offs** - Comprehensiveness vs. token limits
- **Personal style** - There's no single "right" prompt

Important insight: Good prompts are tailored to YOUR specific needs

Tomorrow: Evaluation

The Next Challenge

How Do We Know It's Working?

Today you built prompts and observed outputs.

But:

- How do you evaluate systematically?
- What if you process 50 papers?
- How do you identify failure patterns?
- When is it “good enough”?

Tomorrow (Day 3): We'll tackle systematic evaluation

Homework (Optional)

Continue Experimenting

If you'd like:

- Test your prompt on 2-3 different papers (or other use cases)
- Note what works and what doesn't
- Add observations to the shared doc
- Try refining your prompt further

Question to ponder: What would “good enough” look like for your actual research needs?

Questions & Wrap-Up

Key Takeaways from Day 2

What We Learned

1. **Seven building blocks** provide structure for effective prompts
2. **Iteration is essential** - start simple, add complexity
3. **Specificity matters** - explicit instructions reduce ambiguity
4. **No perfect prompt** - tailor to your needs and domain
5. **AI can help refine** - but understand the principles first

Tomorrow: We'll learn how to evaluate these prompts systematically

Questions?

Before We Close

- Clarifications on any building blocks?
- Challenges with your specific use case?
- Questions about tomorrow?

See you tomorrow for Day 3!