

# What You Need to Know About LLM Training

*Cutting Through the Hype to Understand AI Capabilities*

**Franck Albinet**

*franckalbinet@gmail.com*

*Independent Data Science & AI Consultant*

October 1, 2025

# Why This Matters

# The Challenge

## *Demystifying AI Terminology*

You've heard terms like "reasoning," "intelligence," and "understanding" applied to AI systems.

**The reality:** These terms are a mix of:

- Marketing language to make AI sound impressive
- Engineers needing names for technical processes
- The persistent myth that AI mimics human cognition

# The Challenge (cont.)

## *What Researchers Actually Need*

### **Focus on what matters:**

- Understanding what these systems actually do
- Setting appropriate expectations
- Learning how to interact with them effectively

**Key insight:** AI capabilities are specific, measurable processes—not mysterious intelligence

# The Four Stages of LLM Training

# The Four Stages of LLM Training

## *Our Journey Today*

1. **Self-Supervised Pre-Training** - Building the Foundation
2. **Supervised Fine-Tuning** - Learning to Converse
3. **Preference Fine-Tuning** - Learning What Humans Want
4. **Reasoning Fine-Tuning** - Teaching Step-by-Step Analysis

Each stage builds specific capabilities that inform how we should interact with AI

# The Four Stages of LLM Training

## *Our Journey Today*

1. **Self-Supervised Pre-Training** - Building the Foundation ← *Current*
2. **Supervised Fine-Tuning** - Learning to Converse
3. **Preference Fine-Tuning** - Learning What Humans Want
4. **Reasoning Fine-Tuning** - Teaching Step-by-Step Analysis

# Stage 1: Self-Supervised Pre-Training

## *Building the Foundation*

**The task:** Learn to predict what word comes next

- Model reads massive amounts of text (books, websites, articles)
- Parts are intentionally hidden: "The cat sat on the \_\_\_\_"
- System learns "mat" appears more often than "elephant"

**This is self-supervised learning** - no human labels needed, just pattern recognition from existing text

**What you get:** A system that understands language structure but can't hold conversations



# The Four Stages of LLM Training

## *Our Journey Today*

1. ~~Self-Supervised Pre-Training~~ - Building the Foundation ✓
2. **Supervised Fine-Tuning** - Learning to Converse ← *Current*
3. **Preferred Fine-Tuning** - Learning What Humans Want
4. **Reasoning Fine-Tuning** - Teaching Step-by-Step Analysis

# Stage 2: Supervised Fine-Tuning

## *Learning to Converse*

**The problem:** Text completion  $\neq$  helpful assistant

- Pre-trained model might continue: "This question often comes up..."
- Instead of actually answering your question

**Example:** Ask "What are key success factors?" → gets academic discussion instead of actionable answer

# Stage 2: Supervised Fine-Tuning (cont.)

## *The Solution*

### **Show thousands of human-AI conversations**

- Learn dialogue structure
- Recognize questions and provide direct answers
- Package existing knowledge into useful responses

**What you get:** A conversational system, but it may still produce confident-sounding wrong answers

# The Four Stages of LLM Training

## *Our Journey Today*

1. ~~Self-Supervised Pre-Training~~ - Building the Foundation ✓
2. ~~Supervised Fine-Tuning~~ - Learning to Converse ✓
3. **Preferred Fine-Tuning** - Learning What Humans Want ← *Current*
4. **Reasoning Fine-Tuning** - Teaching Step-by-Step Analysis

# Stage 3: Preference Fine-Tuning

## *Learning What Humans Actually Want*

**The challenge:** How do you train on subjective tasks?

- No single “right” way to write a compelling proposal
- Can’t easily create training examples for complex judgment calls

**Example:** What makes one research summary “better” than another? It’s often a matter of judgment, not facts

# Stage 3: Preference Fine-Tuning (cont.)

## *The Approach*

### **Human feedback on response quality**

- Show model multiple responses to same question
- Learn from human preferences about what's better
- Also handles safety: refusing harmful requests

**What you get:** Responses aligned with human preferences, but may struggle with complex reasoning

# The Four Stages of LLM Training

## *Our Journey Today*

1. ~~Self-Supervised Pre-Training~~ - Building the Foundation ✓
2. ~~Supervised Fine-Tuning~~ - Learning to Converse ✓
3. ~~Preferred Fine-Tuning~~ - Learning What Humans Want ✓
4. **Reasoning Fine-Tuning** - Teaching Step-by-Step Analysis ← *Current*

# Stage 4: Reasoning Fine-Tuning

## *Teaching Step-by-Step Analysis*

**The gap:** Models jump to conclusions without showing work

- Ask for budget evaluation → immediate conclusion
- No systematic breakdown of the problem

**The challenge:** How do we get systematic analysis instead of quick answers?



# Stage 4: Reasoning Fine-Tuning (cont.)

## *The Breakthrough*

**The solution:** Teach models to “think” step-by-step

- Break complex tasks into steps
- Consider multiple approaches
- Show the analytical process

**What you get:** Systems that can tackle multi-step problems and show their work

# Recent Innovations

# Reinforcement Learning

## *The DeepSeek Innovation*

**Traditional approach:** Expensive human evaluation of every response

**DeepSeek's insight:** Focus on verifiable problems

- Math and coding have clear right/wrong answers
- Success can be measured automatically
- Dramatically reduced training costs

**Impact:** Made sophisticated reasoning capabilities more accessible

# The Next Wave: GEPA

*Reflective Prompt Evolution (July 2025)*

**New insight:** Since these are language models, why not teach them using language?

- Natural language feedback instead of numerical rewards
- Models learn to critique their own work
- Self-reflection and improvement through conversation

**Results:** 10-20% better performance with 35× fewer computational resources

# Timeline: How We Got Here

# The Early Foundation (2017-2019)

## *Stage 1: Self-Supervised Pre-Training*

**2017:** Transformer architecture (Google)

- “**Attention is All You Need**” paper revolutionizes NLP

**2018:** First breakthrough models

- GPT-1 (OpenAI) - generative pre-training
- BERT (Google) - bidirectional encoder

**The foundation was set** - but these models couldn't hold conversations

# The Conversation Breakthrough (2020-2021)

## *Stage 2: Supervised Fine-Tuning*

### **2020:** GPT-3 (OpenAI)

- First model that could follow instructions
- Showed emergent abilities at scale

### **2019-2022:** Google's contributions

- T5 (2019) - text-to-text transfer transformer
- PaLM (2022) - 540B parameter breakthrough

**Models could now converse** - but responses weren't always helpful or safe

# The Alignment Revolution (2022)

## *Stage 3: Preference Fine-Tuning*

### **November 2022:** ChatGPT launch (OpenAI)

- RLHF (Reinforcement Learning from Human Feedback)
- Suddenly, everyone had access to AI assistance

### **Rapid competitive response:**

- Claude (Anthropic) - Constitutional AI approach
- Mistral AI emerges with open-source alternatives

The AI race intensifies - every company rushes to catch up



# The Reasoning Era (2024-2025)

## *Stage 4: Reasoning Fine-Tuning*

### **2024:** OpenAI's o1 series

- Chain-of-thought reasoning capabilities
- Models that “think” step-by-step

### **2024-2025:** Efficiency breakthroughs

- DeepSeek-R1 - reasoning at lower cost
- GEPA (July 2025) - 35x more efficient training

**Key insight:** Each breakthrough is adopted and improved by competitors within months, not years

# Implications for Research

# Understanding Capabilities

## *What This Means for Literature Review*

### **Stage 1 (Pre-training):**

- Vast knowledge of scientific literature
- Understanding of domain terminology and concepts

### **Stage 2 (Instruction following):**

- Can respond to specific research queries
- Formats information appropriately

# Understanding Capabilities (cont.)

## *What This Means for Literature Review*

### **Stage 3 (Preference alignment):**

- Produces responses researchers find useful
- Refuses to fabricate citations or make unsupported claims

### **Stage 4 (Reasoning):**

- Can analyze complex research questions systematically
- Shows analytical steps for verification

**Critical insight:** Each stage enables different capabilities—understanding this helps you design better prompts

# Practical Applications

## *How This Informs Your AI Use*

### **Better prompting:**

- Ask for step-by-step analysis (leverages Stage 4)
- Specify output format (leverages Stage 2)
- Request evidence and reasoning (leverages all stages)

**Example:** “Analyze this research question step-by-step and show your reasoning” vs. “What do you think about this?”

# Practical Applications (cont.)

## *Better Evaluation*

### **Check the AI's work systematically:**

- Check if reasoning steps make sense
- Verify factual claims independently
- Look for systematic vs. superficial analysis

**Remember:** Understanding the training helps you work *with* the AI's strengths rather than against its limitations

# Questions & Discussion

# Key Takeaways

## *Moving Forward*

- AI capabilities result from specific training stages
- Each stage builds particular abilities
- Understanding this helps you use AI more effectively
- The field continues evolving rapidly

**Next steps:** How does this change how you think about using AI in your research?



# Further Reading

- **What Decision-Makers Need to Know About AI Training**
- **Constitutional AI**